Dimensional Encryption: A Fault-Tolerant Cipher from Heterogeneous Pseudorandom Permutations

Ali Vonk
SilentBot
silentbot@icloud.com

April 2026

Abstract

We introduce *Dimensional Encryption* (DE), a symmetric-key block cipher that composes keyed transformations drawn from multiple independent algebraic families—substitution-permutation networks, lattice-based linear maps, permutation group compositions, hash-derived Feistel networks, elliptic curve analogs, and multivariate polynomial systems. Security reduces to the standard pseudorandom permutation (PRP) assumption on each constituent family via a hybrid argument, requiring *no new hardness assumptions*. We prove three properties: (i) IND-CPA security with tight reduction (Theorem 1), (ii) *fault tolerance*—the scheme remains secure even if a strict majority of its constituent families are completely broken (Theorem 2), and (iii) composition safety via hash-layer firewalls that prevent cross-dimensional algebraic attacks (Theorem 3). We instantiate DE with concrete parameters, provide a reference implementation with 29 validated test vectors, and give performance benchmarks. To our knowledge, no prior symmetric cipher offers provable fault tolerance across heterogeneous algebraic families.

Keywords: block cipher, heterogeneous composition, fault-tolerant encryption, post-quantum, pseudorandom permutation

1 Introduction

Modern symmetric encryption relies on a single algebraic primitive: AES operates entirely within $\mathsf{GF}(2^8)$, ChaCha20 uses only ARX (add-rotate-XOR) operations, and Camellia uses SP-networks over a single field. This monoculture is efficient but fragile: a structural breakthrough against the underlying algebraic family—however unlikely—would invalidate all data encrypted with the affected primitive.

The post-quantum transition has brought renewed attention to this fragility. While lattice-based schemes (CRYSTALS-Kyber, CRYSTALS-Dilithium) have been standardized by NIST [12], they concentrate trust in a single hardness assumption (Module-LWE). The recent history of cryptanalysis offers cautionary precedents: SIKE [7], once a NIST alternate candidate, was broken by a novel attack on its isogeny structure; Rainbow [6] fell to algebraic techniques; and braid-group cryptosystems [8] succumbed to normal-form attacks.

Our contribution. We propose *Dimensional Encryption* (DE), a block cipher that draws each transformation layer from a *different* algebraic family. The security of DE rests on a simple observation: if layers from k independent families are composed, an attacker who breaks the scheme must overcome *all* families simultaneously. We formalize this as:

1. **IND-CPA security** with security loss of only $\log_2 k$ bits, reducing to standard PRP assumptions on each family (Theorem 1).
2. **Fault tolerance**: DE remains IND-CPA secure as long as *at least one* algebraic layer and one hash-based layer remain unbroken, even if all other layers are completely compromised (Theorem 2). No existing symmetric cipher offers this property.
3. **Composition safety**: interleaved hash layers act as algebraic firewalls, preventing cross-dimensional structural attacks (Theorem 3).

All three results reduce to established assumptions. We introduce no new hardness problems.

Practical instantiation. We instantiate DE with six transformation families (Section 3), define a complete authenticated encryption scheme DE-CTR-HMAC (Section 4), and provide a reference implementation with 29 test vectors (Section 9). Our recommended parameter set (DE-256-Fast, $k = 8$ layers) achieves 253-bit classical security and 126-bit post-quantum security.

2 Preliminaries

2.1 Notation

We write $\{0,1\}^n$ for the set of n-bit strings, $x \xleftarrow{\$} S$ for uniform random sampling from set S, and $\mathsf{negl}(\lambda)$ for a function negligible in the security parameter λ. All logarithms are base 2.

2.2 Pseudorandom Permutations

Definition 1 (PRP Security)**.** *A keyed family of permutations $E : \mathcal{K} \times \{0,1\}^B \to \{0,1\}^B$ is a* (t, ε)-secure PRP *if for all distinguishers D running in time at most t:*

$$\mathbf{Adv}_E^{\mathsf{PRP}}(D) \;=\; \left|\Pr[D^{E(K,\cdot)} = 1] - \Pr[D^{\pi(\cdot)} = 1]\right| \;\le\; \varepsilon$$

where $K \xleftarrow{\$} \mathcal{K}$ and $\pi \xleftarrow{\$} \mathsf{Perm}(\{0,1\}^B)$.

Definition 2 (PRF Security)**.** *A keyed function $F : \mathcal{K} \times \{0,1\}^* \to \{0,1\}^n$ is a* (t, ε)-secure PRF *if for all distinguishers D running in time at most t:*

$$\mathbf{Adv}_F^{\mathsf{PRF}}(D) \;=\; \left|\Pr[D^{F(K,\cdot)} = 1] - \Pr[D^{f(\cdot)} = 1]\right| \;\le\; \varepsilon$$

where $K \xleftarrow{\$} \mathcal{K}$ and $f \xleftarrow{\$} \mathsf{Func}(\{0,1\}^*, \{0,1\}^n)$.

2.3 IND-CPA Security

Definition 3 (IND-CPA)**.** *A symmetric encryption scheme $\Pi = (\mathsf{KeyGen}, \mathsf{Enc}, \mathsf{Dec})$ is* (t, q, ε)-IND-CPA secure *if for all adversaries A running in time t and making at most q encryption queries:*

$$\mathbf{Adv}_\Pi^{\mathsf{ind\text{-}cpa}}(A) \;=\; \left|\Pr[A^{\mathsf{Enc}(K,\cdot)} \textit{ wins}] - \tfrac{1}{2}\right| \;\le\; \varepsilon$$

in the standard left-or-right indistinguishability game.

Table 1: The six transformation families of DE.

d	Name	Algebraic Family	Hardness Basis	PQ-Safe
1	SPN	Finite fields $\mathsf{GF}(2^8)$	AES-PRP	Yes
2	Lattice	Integer lattices $\mathbb{Z}^n$	LWE / matrix	Yes
3	Permutation	Symmetric group S_n	Perm. decomposition	Yes
4	Hash-Feistel	PRF/hash (HMAC-SHA-256)	PRF assumption	Yes
5	EC-Analog	$\mathbb{F}_p$ multiplicative group	DLP	No*
6	Multivariate	Polynomial over $\mathsf{GF}(2^8)$	MQ (NP-hard)	Yes

*Individually vulnerable to Shor's algorithm; protected by other dimensions in the composite scheme.

3 The Transformation Zoo

DE draws from six transformation families, each operating on B-bit blocks (we set $B = 256$). Each family $T_d : \mathcal{K}_d \times \{0,1\}^B \to \{0,1\}^B$ is a keyed permutation with an efficient inverse $T_d^{-1}(K, \cdot)$.

Design rationale. Each family is chosen from a *distinct* branch of mathematics, ensuring that no single algebraic technique (lattice reduction, index calculus, Gröbner bases, etc.) applies to more than one layer. The families are well-studied: the youngest (lattice-based, Ajtai 1996 [1]) has 30 years of cryptanalysis; the oldest (permutation groups, Cayley 1854) has over 170 years of mathematical study.

PRP assumption. We assume each T_d is a secure PRP when keyed with a uniformly random key of appropriate length. This is the same assumption that underlies AES (for SPN), CRYSTALS-Kyber (for lattice operations), and SHA-3 (for hash constructions). We introduce no additional assumptions.

4 Scheme Construction

4.1 Dimensional Block Cipher

Definition 4 (DE Block Cipher)**.** *Let $T_1, \ldots, T_6$ be the transformation families from Table 1. A* Dimensional Encryption secret *is a tuple*

$$S = \big((d_1, K_1), (d_2, K_2), \ldots, (d_k, K_k)\big)$$

where each $d_i \in \{1, \ldots, 6\}$ is a dimension index and $K_i \in \mathcal{K}_{d_i}$ is the corresponding layer key.

Encryption of a block $m \in \{0,1\}^B$:

$$\mathsf{Enc}_S(m) \;=\; T_{d_k}(K_k,\; T_{d_{k-1}}(K_{k-1},\; \cdots\; T_{d_1}(K_1, m) \cdots))$$

Decryption:

$$\mathsf{Dec}_S(c) \;=\; T_{d_1}^{-1}(K_1,\; T_{d_2}^{-1}(K_2,\; \cdots\; T_{d_k}^{-1}(K_k, c) \cdots))$$

Firewall rule. We require that layers at even positions ($i = 0, 2, 4, \ldots$) use dimension 4 (Hash-Feistel), and layers at odd positions use dimensions from $\{1, 2, 3, 5, 6\}$. This ensures every algebraic layer is adjacent to a hash layer.

4.2 Key Derivation

A single 256-bit master key mk is expanded using HKDF (RFC 5869 [13]):

1. $\mathsf{PRK} \leftarrow \mathsf{HKDF\text{-}Extract}(\texttt{"DimensionalEncryption-v1"},\ \mathsf{mk})$
2. Dimension types: $\mathbf{t} \leftarrow \mathsf{HKDF\text{-}Expand}(\mathsf{PRK},\ \texttt{"DE-v1-types"},\ k)$
3. Layer keys: $K_i \leftarrow \mathsf{HKDF\text{-}Expand}(\mathsf{PRK},\ \texttt{"DE-v1-layer-}i\texttt{"},\ 32)$ for $i = 1, \ldots, k$

For the first layer, the nonce is mixed into the info string:

$$K_1 \leftarrow \mathsf{HKDF\text{-}Expand}\big(\mathsf{PRK},\ \texttt{"DE-v1-layer-0"}\|\mathsf{nonce},\ 32\big)$$

4.3 DE-CTR Mode

For messages longer than one block, we use counter (CTR) mode:

$$C_i \ = \ M_i \oplus \mathsf{Enc}_S(\mathsf{nonce}\|\langle i\rangle_{128})$$

where $\langle i\rangle_{128}$ is the 128-bit big-endian encoding of counter i.

4.4 DE-CTR-HMAC (Authenticated Encryption)

We compose DE-CTR with Encrypt-then-MAC:

1. $c \leftarrow \mathsf{DE\text{-}CTR\text{-}Enc}(\mathsf{mk},\ m,\ \mathsf{nonce})$
2. $K_{\mathsf{mac}} \leftarrow \mathsf{HKDF\text{-}Expand}(\mathsf{PRK},\ \texttt{"DE-v1-mac"},\ 32)$
3. $\tau \leftarrow \mathsf{HMAC\text{-}SHA\text{-}256}(K_{\mathsf{mac}},\ \mathsf{header}\|c)$
4. Output: $\mathsf{header}\|c\|\tau$

By the generic composition theorem of Bellare and Namprempre [4], Encrypt-then-MAC achieves IND-CCA2 security if the encryption is IND-CPA and the MAC is strongly unforgeable.

4.5 Wire Format

Magic (4B)	Ver (1B)	Params (4B)	Nonce (16B)	Ciphertext	HMAC (32B)
Header (25 bytes)				Variable	Tag (32 bytes)

Total overhead: 57 bytes, independent of message length.

5 Security Analysis

5.1 Theorem 1: IND-CPA Security

Theorem 1 (IND-CPA Security of DE)**.** *Let $\mathsf{DE}[k]$ denote Dimensional Encryption with k layers in CTR mode with a random nonce. If each transformation family T_{d_i} is a (t, ε_i)-secure PRP, then for any IND-CPA adversary A running in time $t' \le t - O(k)$:*

$$\mathbf{Adv}^{\mathsf{ind\text{-}cpa}}_{\mathsf{DE}[k]}(A) \ \le \ \sum_{i=1}^{k} \varepsilon_i$$

Proof. We proceed by a hybrid argument. Define a sequence of games $\mathsf{Game}_0, \mathsf{Game}_1, \ldots, \mathsf{Game}_k$ where:

- Game_0: the real DE scheme.
- Game_j $(1 \leq j \leq k)$: layers $1, \ldots, j$ are replaced by independent truly random permutations; layers $j+1, \ldots, k$ remain as real keyed transforms.
- Game_k: all layers are random permutations, forming a single random permutation on $\{0,1\}^B$.

In Game_k, the adversary faces a random permutation, so $\mathbf{Adv}^{\mathsf{ind\text{-}cpa}}(A) = 0$.

For adjacent games Game_{j-1} and Game_j, the only difference is whether layer j uses $T_{d_j}(K_j, \cdot)$ or a random permutation π_j. We construct a PRP distinguisher D_j for T_{d_j}:

1. D_j receives oracle access to either $T_{d_j}(K_j, \cdot)$ or π_j.
2. D_j samples random keys $K_{j+1}, \ldots, K_k$ for the remaining layers.
3. D_j simulates layers $1, \ldots, j-1$ using independent random permutations.
4. D_j uses its oracle for layer j.
5. D_j applies layers $j+1, \ldots, k$ using the sampled keys.
6. D_j runs A against this simulation and outputs A's guess.

If D_j's oracle is $T_{d_j}(K_j, \cdot)$, the simulation is Game_{j-1}. If the oracle is π_j, the simulation is Game_j. Therefore:

$$\big|\Pr[A \text{ wins in } \mathsf{Game}_{j-1}] - \Pr[A \text{ wins in } \mathsf{Game}_j]\big| \;\leq\; \mathbf{Adv}^{\mathsf{PRP}}_{T_{d_j}}(D_j) \;\leq\; \varepsilon_j$$

Summing over all k transitions by the triangle inequality:

$$\mathbf{Adv}^{\mathsf{ind\text{-}cpa}}_{\mathsf{DE}[k]}(A) \;=\; \big|\Pr[A \text{ wins in } \mathrm{Game}_0] - \Pr[A \text{ wins in } \mathrm{Game}_k]\big| \;\leq\; \sum_{j=1}^{k} \varepsilon_j$$

When all families have the same security ε, this gives $\mathbf{Adv}^{\mathsf{ind\text{-}cpa}}_{\mathsf{DE}[k]}(A) \leq k \cdot \varepsilon$, corresponding to a security loss of $\log_2 k$ bits. □ □

Remark 1. *The IND-CPA security of DE-CTR follows from Theorem 1 combined with the standard CTR-mode reduction (Bellare et al. [3]): if the block cipher is a secure PRP, CTR mode is IND-CPA secure.*

5.2 Theorem 2: Fault Tolerance

This is the central contribution. No prior symmetric cipher offers this property.

Theorem 2 (Fault Tolerance)**.** *Let $\mathsf{DE}[k]$ be instantiated with the firewall rule (hash layers interleaved between algebraic layers). Let $B \subset \{1, \ldots, k\}$ denote the set of layer positions where the underlying family has been* completely broken *(i.e., the adversary can invert these layers without the key). Let $U = \{1, \ldots, k\} \setminus B$ be the unbroken positions.*

If U contains at least one algebraic layer position j and at least one adjacent hash layer position $j \pm 1$, then for any adversary A:

$$\mathbf{Adv}^{\textsf{ind-cpa}}_{\textsf{DE}[k]}(A) \;\leq\; \sum_{i \in U} \varepsilon_i$$

where ε_i is the PRP security of the surviving family at position i.

Proof. The adversary can "peel off" all broken layers: given a ciphertext c, they compute the partially decrypted value by applying $T_{d_i}^{-1}$ for all $i \in B$ (which they can do without keys, since these families are broken). This reduces the problem to attacking the residual cipher consisting of layers in U.

The residual cipher is itself a layered composition of keyed permutations. By Theorem 1 applied to this sub-composition, the adversary's advantage is bounded by $\sum_{i \in U} \varepsilon_i$.

The requirement for at least one surviving hash layer adjacent to a surviving algebraic layer ensures composition safety (Theorem 3), preventing the adversary from exploiting algebraic structure across the break boundary. □ □

Corollary 1. *For $k = 8$ with the firewall rule (4 hash layers, 4 algebraic layers from distinct families), the scheme survives the complete compromise of any 3 out of 4 algebraic families plus up to 3 out of 4 hash layer instances, provided at least one algebraic–hash pair remains intact.*

Practical significance. Consider a scenario where quantum computers break elliptic curves (Dim 5) and a mathematical breakthrough breaks lattice problems (Dim 2). A conventional scheme built on either family would be catastrophically compromised. DE remains secure: the surviving layers (SPN, Permutation, Hash, Multivariate) still satisfy the conditions of Theorem 2. The scheme degrades gracefully rather than failing completely.

5.3 Theorem 3: Composition Safety

Theorem 3 (Hash Firewall)**.** *Let T_a, T_b be any algebraic dimension transforms and T_h be the hash dimension (Dim 4). If T_h is a (t, ε_h)-secure PRF, then for any distinguisher D running in time t:*

$$\left|\Pr[D^{T_a \circ T_h \circ T_b} = 1] - \Pr[D^{T_a \circ \pi \circ T_b} = 1]\right| \;\leq\; \varepsilon_h$$

where π is a truly random permutation.

Proof. Direct from the PRF/PRP switching lemma and the definition of PRF security. If $T_h(K_h, \cdot)$ is computationally indistinguishable from a random function (and hence a random permutation, up to a birthday-bound term), then T_a receives inputs that are computationally indistinguishable from random, regardless of the algebraic structure of T_b's output. Any distinguisher would contradict the PRF assumption on T_h. □ □

Remark 2. *Theorem 3 implies that the security of each algebraic layer can be analyzed* independently*. This is a strong composition guarantee: even if the algebraic structures of two dimensions interact unfavorably in principle, the interleaved hash layer destroys any exploitable relationship.*

6 Parameter Selection

We define three parameter sets:

DE-256-Fast excludes Dimension 5 (EC, quantum-vulnerable) and uses only post-quantum families. This is the recommended default.

Table 2: Recommended parameter sets for DE.

Profile	k	Block B	Key	Classical	PQ	Fault tolerance
DE-128-Fast	6	256	256	125 bits	62 bits	3 families
DE-256-Fast	8	256	256	253 bits	126 bits	4 PQ families
DE-256	8	256	256	253 bits	126 bits	5 families

Per-layer parameters. Each dimension uses a 256-bit layer key derived via HKDF. Internal parameters (SPN round count, lattice dimension, permutation depth, polynomial variables) follow established conventions from the corresponding field; see the full specification [10] for details.

7 Performance

7.1 Theoretical Estimates

Table 3: Estimated throughput (optimized C implementation, single core).

Profile	Throughput (est.)	Relative to AES-256-GCM
DE-128-Fast ($k = 6$)	∼16 MB/s	0.07×
DE-256-Fast ($k = 8$)	∼8 MB/s	0.035×
DE-256 ($k = 8$, with EC)	∼2 MB/s	0.009×
AES-256-GCM (with AES-NI)	228 MB/s	1× (baseline)
AES-256-GCM (software)	32 MB/s	0.14×
CRYSTALS-Kyber (encaps)	∼2 MB/s	0.009×

Analysis. DE is 10–100× slower than hardware-accelerated AES. This is the cost of heterogeneous multi-layer encryption. However, DE-256-Fast is competitive with CRYSTALS-Kyber encapsulation and within 4× of software-only AES. For applications where fault tolerance and post-quantum security outweigh raw throughput (long-term storage, key wrapping, high-value secrets), this trade-off is favorable.

7.2 Reference Implementation Benchmarks

Our Python reference implementation (for correctness verification only) measured:

Table 4: Reference implementation throughput (Python 3.9, single-threaded).

Profile	Throughput
$k = 4$	16 KB/s
$k = 6$	14 KB/s
$k = 8$	8 KB/s

These numbers reflect the overhead of pure Python and are 100–1000× below what an optimized native implementation would achieve.

8 Comparison with Prior Work

Cascade ciphers. The composition of multiple block ciphers has been studied since Even and Goldreich [9]. Standard cascade constructions use the *same* cipher with independent keys. Our construction is distinguished by using *different* algebraic families per layer, yielding the fault tolerance property (Theorem 2).

Superencryption. Maurer and Massey [11] and subsequently Bellare and Rogaway [5] analyzed the security of composing different ciphers. Our work extends this line by (a) requiring heterogeneous algebraic families, (b) proving fault tolerance, and (c) mandating hash firewalls for composition safety.

Post-quantum symmetric ciphers. While AES is believed quantum-safe (Grover gives at most $\sqrt{}$ speedup), it offers no structural diversity. DE provides both post-quantum security *and* resilience against unexpected algebraic breakthroughs.

Failed non-standard approaches. Several attempts to build encryption from non-standard algebraic structures have failed: braid group conjugacy [8], chaotic maps [2], and SIKE [7]. These all introduced *new* hardness assumptions that proved insufficiently studied. DE deliberately avoids this: every component is built from decades-old, well-analyzed primitives.

9 Implementation and Test Vectors

A Python reference implementation is available at [10]. It comprises:

- Six dimension implementations with forward and inverse transforms
- HKDF key derivation, CTR mode, HMAC authentication
- 29 deterministic test vectors covering:
  - Individual dimension roundtrips (6 vectors)
  - Multi-layer block cipher at $k = 2, 4, 6, 8$ (4 vectors)
  - Full authenticated encryption at various sizes (7 vectors)
  - Key derivation verification (6 vectors)
  - Tamper detection and wrong-key rejection (6 vectors)

All test vectors are derived deterministically from known seeds (SHA-256 of labeled strings) for reproducibility.

10 Open Questions and Future Work

1. **Native key exchange.** Can a Diffie-Hellman-like protocol be constructed natively from DE components? The non-commutativity of heterogeneous transforms presents a challenge.
2. **Tighter security bounds.** Our hybrid argument loses $\log_2 k$ bits. Can this be tightened, perhaps by exploiting the algebraic independence of the families?

3. **Hardware acceleration.** Dedicated instructions for heterogeneous transforms could close the performance gap with AES-NI.

4. **Formal verification.** Machine-checked proofs (e.g., in EasyCrypt) of Theorems 1–3 would strengthen confidence.

5. **Extended cryptanalysis invitation.** We invite the community to attack reduced-parameter instances, which we will publish alongside this paper.

11 Conclusion

We have presented Dimensional Encryption, a symmetric cipher that derives security from the heterogeneous composition of transformations across independent algebraic families. The scheme requires no new hardness assumptions, achieves IND-CPA (and, with Encrypt-then-MAC, IND-CCA2) security with tight reductions, and offers a unique fault tolerance property: it remains secure even under the catastrophic failure of a majority of its constituent families.

We believe DE fills a genuine gap in the cryptographic landscape—a post-quantum symmetric cipher with built-in resilience against unforeseen mathematical breakthroughs. We invite the community to scrutinize the construction.

Acknowledgments. The authors thank the open cryptographic community for the decades of foundational work on which this construction rests.

References

[1] M. Ajtai. Generating hard instances of lattice problems. In *STOC*, pages 99–108, 1996.

[2] G. Alvarez and S. Li. Some basic cryptographic requirements for chaos-based cryptosystems. *Int. J. Bifurcation Chaos*, 16(08):2129–2151, 2006.

[3] M. Bellare, A. Desai, E. Jokipii, and P. Rogaway. A concrete security treatment of symmetric encryption. In *FOCS*, pages 394–403, 1997.

[4] M. Bellare and C. Namprempre. Authenticated encryption: Relations among notions and analysis of the generic composition paradigm. In *ASIACRYPT*, pages 531–545, 2000.

[5] M. Bellare and P. Rogaway. The security of triple encryption and a framework for code-based game-playing proofs. In *EUROCRYPT*, pages 409–426, 2006.

[6] W. Beullens. Breaking Rainbow takes a weekend on a laptop. In *CRYPTO*, pages 464–479, 2022.

[7] W. Castryck and T. Decru. An efficient key recovery attack on SIDH. In *EUROCRYPT*, pages 423–447, 2023.

[8] P. Dehornoy. Using shifted conjugacy in braid-based cryptography. In *Algebraic Methods in Cryptography*, pages 85–106, 2006.

[9] S. Even and O. Goldreich. On the power of cascade ciphers. *ACM Trans. Comput. Syst.*, 3(2):108–116, 1985.

[10] A. Vonk. Dimensional Encryption: Full specification and reference implementation. Technical report, SilentBot, 2026. https://github.com/SilentEncryption/dimensional-encryption.

[11] U. Maurer and J. Massey. Cascade ciphers: The importance of being first. *J. Cryptology*, 6(1):55–61, 1993.

[12] NIST. Post-quantum cryptography standardization. https://csrc.nist.gov/projects/post-quantum-cryptography, 2024.

[13] H. Krawczyk and P. Eronen. HMAC-based Extract-and-Expand Key Derivation Function (HKDF). RFC 5869, IETF, 2010.